

Aim

To perform customer segmentation using K-Means clustering on the Online Retail dataset and identify outliers using various methods.

Algorithm

1. Load and preprocess the dataset (handle missing values, filter out invalid entries).
2. Encode categorical variables and scale numerical features.
3. Determine the optimal number of clusters (K) using the Elbow method and Silhouette score.
4. Apply K-Means clustering and assign labels to data points.
5. Detect outliers using distance from cluster centers, Z-score method, and Isolation Forest.
6. Visualize clustering results using PCA.

Algorithm Description

K-Means is an unsupervised clustering algorithm that partitions data into K clusters by minimizing intra-cluster variance. The algorithm iteratively:

- Assigns each data point to the nearest cluster center.
- Updates cluster centers by computing the mean of assigned points.
- Repeats until convergence.

Outlier detection methods include:

- **Distance-based outliers:** Identifies data points far from cluster centers.
- **Z-Score method:** Flags points with values beyond three standard deviations.
- **Isolation Forest:** Detects anomalies based on recursive partitioning of data.

Results

- The optimal number of clusters (K) was determined using the silhouette score.
- Clustering results were visualized using PCA-reduced data.
- Outliers were successfully detected using multiple methods.
- The model provides customer segmentation insights based on purchase behavior.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score, davies_bouldin_score
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import pairwise_distances_argmin_min
from scipy.stats import zscore
from sklearn.ensemble import IsolationForest
```

```
df = pd.read_excel('Online Retail.xlsx')
df
```

	InvoiceNo	StockCode	Description
Quantity \			
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER
6			
1	536365	71053	WHITE METAL LANTERN
6			
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER
8			
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE
6			
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.
6			
...
...			
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS
12			
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL
6			
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL
4			
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE
4			
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT
3			

	InvoiceDate	UnitPrice	CustomerID	Country
0	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
...
541904	2011-12-09 12:50:00	0.85	12680.0	France
541905	2011-12-09 12:50:00	2.10	12680.0	France
541906	2011-12-09 12:50:00	4.15	12680.0	France
541907	2011-12-09 12:50:00	4.15	12680.0	France
541908	2011-12-09 12:50:00	4.95	12680.0	France

[541909 rows x 8 columns]

df.shape

(541909, 8)

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 541909 entries, 0 to 541908

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	InvoiceNo	541909 non-null	object
1	StockCode	541909 non-null	object
2	Description	540455 non-null	object
3	Quantity	541909 non-null	int64
4	InvoiceDate	541909 non-null	datetime64[ns]
5	UnitPrice	541909 non-null	float64
6	CustomerID	406829 non-null	float64
7	Country	541909 non-null	object

dtypes: datetime64[ns](1), float64(2), int64(1), object(4)

memory usage: 33.1+ MB

df.isna().sum()

InvoiceNo	0
StockCode	0
Description	1454
Quantity	0
InvoiceDate	0
UnitPrice	0
CustomerID	135080
Country	0

dtype: int64

```
df['InvoiceDate'] = df['InvoiceDate'].apply(lambda x: str(x).split()[0].split('-')[0])
df
```

	InvoiceNo	StockCode	Description
Quantity \			
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER
6			
1	536365	71053	WHITE METAL LANTERN
6			
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER
8			
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE
6			
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.
6			
...
...			
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS
12			
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL
6			
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL
4			
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE
4			
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT
3			

	InvoiceDate	UnitPrice	CustomerID	Country
0	2010	2.55	17850.0	United Kingdom
1	2010	3.39	17850.0	United Kingdom
2	2010	2.75	17850.0	United Kingdom
3	2010	3.39	17850.0	United Kingdom
4	2010	3.39	17850.0	United Kingdom
...
541904	2011	0.85	12680.0	France
541905	2011	2.10	12680.0	France
541906	2011	4.15	12680.0	France
541907	2011	4.15	12680.0	France
541908	2011	4.95	12680.0	France

[541909 rows x 8 columns]

```
df = df.dropna(subset=['Description'])
df.shape
```

(540455, 8)

```
print(f"Missing values after cleaning: {df.isna().sum().sum()}")
```

Missing values after cleaning: 133626

```
imputer_num = SimpleImputer(strategy='median')  
df.loc[:, 'CustomerID'] =  
imputer_num.fit_transform(df[['CustomerID']])
```

```
df = df[df['Quantity'] > 0]  
df = df[df['UnitPrice'] > 0]
```

```
df.shape
```

```
(530104, 8)
```

```
df.isna().sum().sum()
```

```
0
```

```
label_encoder = LabelEncoder()
```

```
for col in df.columns:
```

```
    if df[col].dtype == 'object':  
        df[col] = df[col].astype(str)  
        df[col] = label_encoder.fit_transform(df[col])
```

```
df
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate
UnitPrice \					
0	0	3407	3844	6	0
2.55					
1	0	2729	3852	6	0
3.39					
2	0	2953	888	8	0
2.75					
3	0	2897	1859	6	0
3.39					
4	0	2896	2849	6	0
3.39					
...
...					
541904	19958	1489	2321	12	1
0.85					
541905	19958	1765	718	6	1
2.10					
541906	19958	2105	724	4	1
4.15					
541907	19958	2106	723	4	1
4.15					
541908	19958	1056	282	3	1
4.95					

	CustomerID	Country
0	17850.0	36
1	17850.0	36
2	17850.0	36
3	17850.0	36
4	17850.0	36
...
541904	12680.0	13
541905	12680.0	13
541906	12680.0	13
541907	12680.0	13
541908	12680.0	13

[530104 rows x 8 columns]

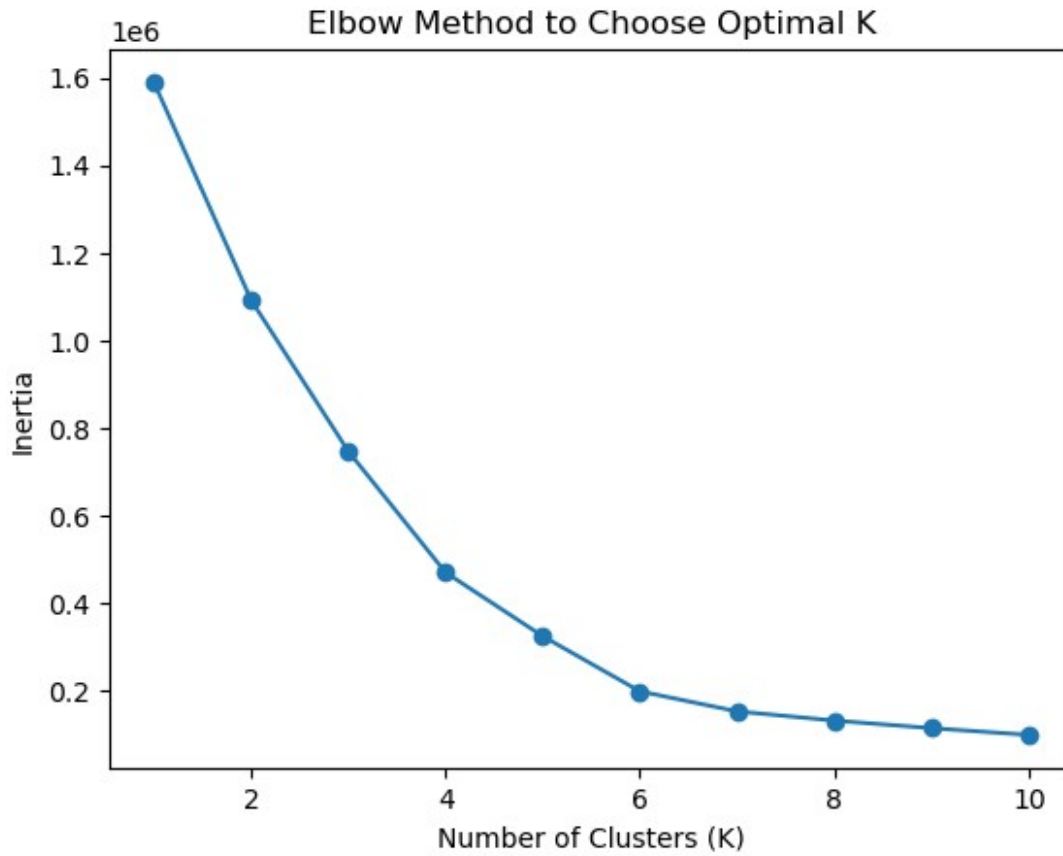
```

features = df[['Quantity', 'UnitPrice', 'CustomerID']]
scaler = StandardScaler()
scaled_data = scaler.fit_transform(features)

inertia = []
k_range = range(1, 11)
for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(scaled_data)
    inertia.append(kmeans.inertia_)

plt.plot(k_range, inertia, marker='o')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Inertia')
plt.title('Elbow Method to Choose Optimal K')
plt.show()

```



```
silhouette_scores = []
for k in k_range[1:]:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
    kmeans.fit(scaled_data)
    score = silhouette_score(scaled_data, kmeans.labels_)
    silhouette_scores.append(score)

plt.plot(k_range[1:], silhouette_scores, marker='o')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Silhouette Score')
plt.title('Silhouette Score for Different K Values')
plt.show()

optimal_k = k_range[1:][np.argmax(silhouette_scores)]
print(f"Optimal number of clusters (K) based on silhouette score:
{optimal_k}")

param_grid = {
    'n_clusters': [optimal_k, 4, 6],
    'n_init': [10],
    'init': ['k-means++']
}
```

```

best_silhouette_score = -1
best_kmeans = None
best_params = None

for n_clusters in param_grid['n_clusters']:
    for init in param_grid['init']:
        kmeans = KMeans(n_clusters=n_clusters,
n_init=param_grid['n_init'][0], init=init, random_state=42)
        kmeans.fit(scaled_data)
        score = silhouette_score(scaled_data, kmeans.labels_)
        if score > best_silhouette_score:
            best_silhouette_score = score
            best_kmeans = kmeans
            best_params = {
                'n_clusters': n_clusters,
                'init': init
            }

df['Cluster'] = best_kmeans.labels_

distances = pairwise_distances_argmin_min(scaled_data,
best_kmeans.cluster_centers_)[1]

threshold = np.percentile(distances, 95)
df['Outlier'] = np.where(distances > threshold, 1, 0)

pca = PCA(n_components=2)
pca_components = pca.fit_transform(scaled_data)

plt.figure(figsize=(8, 6))
plt.scatter(pca_components[:, 0], pca_components[:, 1],
c=df['Cluster'], cmap='viridis', marker='o', label='Clusters')
plt.scatter(pca_components[df['Outlier'] == 1, 0],
pca_components[df['Outlier'] == 1, 1], color='red', label='Outliers',
marker='x')
plt.title(f'K-means Clustering with K={optimal_k} and Outliers')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.legend()
plt.show()

print(f"Cluster Centers:\n{kmeans.cluster_centers_}")

df['Z_Score_Quantity'] = np.abs(zscore(df['Quantity']))
df['Z_Score_Price'] = np.abs(zscore(df['UnitPrice']))
df['Outlier_ZScore'] = np.where((df['Z_Score_Quantity'] > 3) |
(df['Z_Score_Price'] > 3), 1, 0)

iso_forest = IsolationForest(contamination=0.05, random_state=42)
df['Outlier_IsoForest'] = iso_forest.fit_predict(scaled_data)

```



```
df['Outlier_IsoForest'] = df['Outlier_IsoForest'].apply(lambda x: 1 if
x == -1 else 0)

outliers_tuned_df = df[df['Outlier'] == 1]
outliers_zscore_df = df[df['Outlier_ZScore'] == 1]
outliers_iso_df = df[df['Outlier_IsoForest'] == 1]

print("Outliers detected after KMeans:")
print(outliers_tuned_df)
print("\nOutliers detected using Z-Score method:")
print(outliers_zscore_df)
print("\nOutliers detected using Isolation Forest:")
print(outliers_iso_df)

print(f"Cluster Centers:\n{best_kmeans.cluster_centers_}")
```